


RAJU GANDHI

---

# DEEP DIVE INTO DOCKERFILES



# RAJU GANDHI

   @LOOSELYTYPED

CTO - INTEGRALLIS SOFTWARE

# DOCKERFILES?

# THE WHAT AND WHY

- ▶ A set of instructions to build a Docker image
- ▶ Plain text, version controlled
- ▶ Provides insight into the image needs/capabilities/intents

# UNION FILE SYSTEM

bd03254d7d98

FROM looselytyped/java-on-ubuntu

4d3d9ad31e2f

ARG VERSION

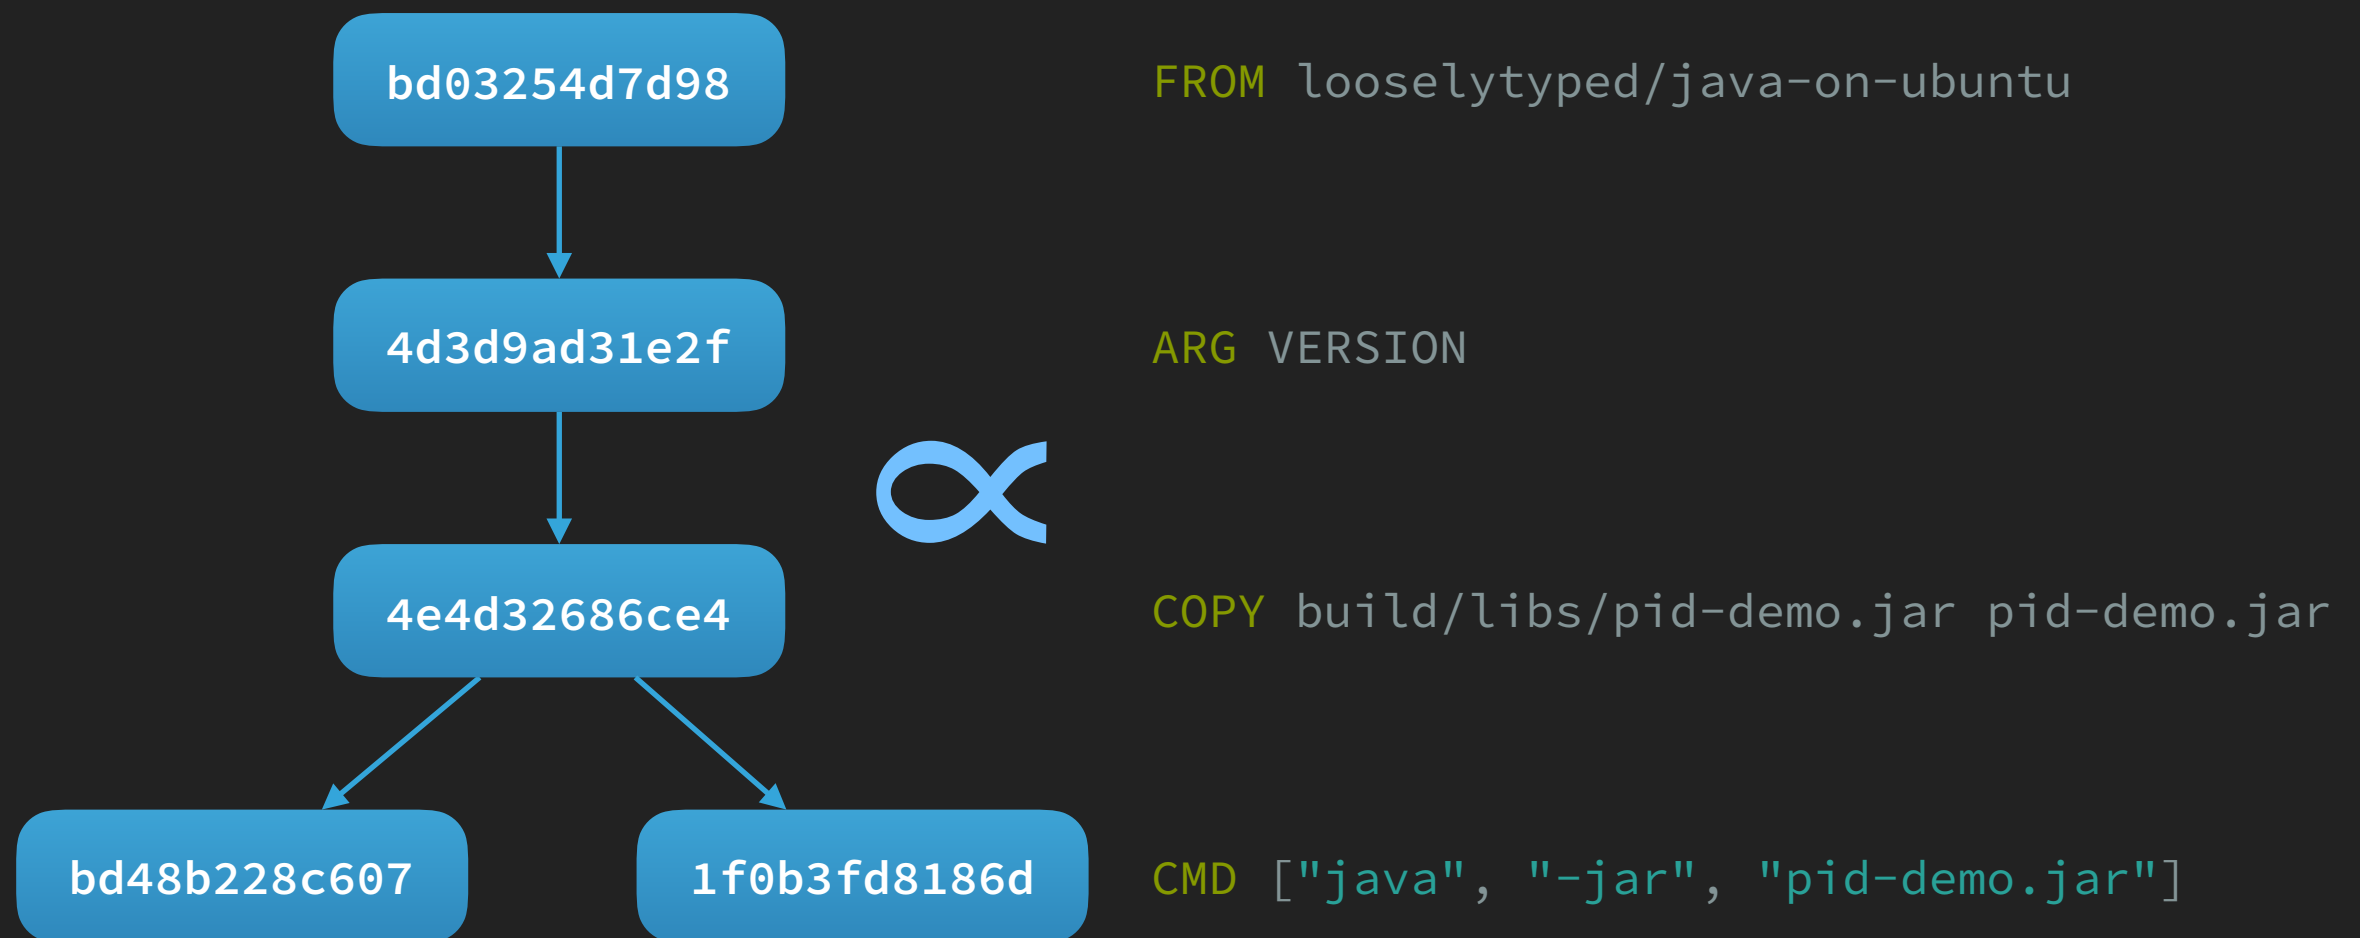
4e4d32686ce4

COPY build/libs/pid-demo.jar pid-demo.jar

bd48b228c607

CMD ["java", "-jar", "pid-demo.jar"]







### THE CODEZ

- ▶ Each instruction adds a “write-only” layer to the UFS
- ▶ Layers are cached for future builds
- ▶ Deletes *increase* the image size

# DON'TS



- ▶ Don't sacrifice readability/maintainability!
- ▶ Don't add unnecessary instructions
- ▶ Beware that whitespace changes can invalidate cache
  - ▶ Think of formatting upfront

### DO'S



- ▶ Be cognizant of layers and image sizes
- ▶ Order matters!
  - ▶ Keep the “moving” bits lower in the Dockerfile
  - ▶ Keep “informational” bits lower in the Dockerfile

**ONE CONCERN PER  
CONTAINER ... FOR THE MOST  
PART**

**@LOOSELYTYPED**

FROM

# NOTES

- ▶ Implies "ancestry"
- ▶ Has to be the first line
- ▶ Has implications on WORKDIR, USER, ENTRYPOINT (and CMD), and ONBUILD, EXPOSE and other commands
- ▶ Create a base image with FROM scratch

FROM

---

## DON'TS



- ▶ Multiple FROMs are allowed. Do **NOT** do [this](#)

## DO'S



- ▶ Pin down the exact tag
  - ▶ Do not use "latest" tag
- ▶ Inspect ancestor images for USERS, PORTs, ENVs, VOLUMEs, LABELs and anything that can be inherited



**RUN**

## DON'TS



- ▶ Be cognizant of the effects (and drawbacks) of caching
- ▶ Do not do OS level upgrades (eg. `RUN dist-upgrade`)

## DOS



- ▶ Group common operations
  - ▶ Clean up as well (reduces image sizes)
- ▶ Use multiline (\) to make PR / auditing easier

**ADD/COPY**

## DON'TS



- ▶ Avoid ADD
- ▶ Do not leave “residual” artifacts

### DO'S



- ▶ Instead of ADD
  - ▶ Combine COPY and RUN
  - ▶ OR RUN with wget/curl
  - ▶ See DO'S under RUN
- ▶ Be mindful of what you put in the `.dockerignore` file

# USER

# DON'TS



- ▶ Do not switch USER often
- ▶ Avoid using root



## DO'S



- ▶ Create a user (if you can) for your service
  - ▶ **RUN** `groupadd -r myuser && useradd -r -g myapp myuser`
  - ▶ `USER myuser`

# **LABELS**

### DONT'S



- ▶ Define individual labels separately

### DO'S



- ▶ Use them liberally
- ▶ Labels can see ENV variables. Use this!
  - ▶ BUILD\_NUMBER, GIT\_SHA
- ▶ Use both image (compile) and container (run) LABELS
- ▶ Apply a standard [convention](#)
- ▶ Build tooling on top of the conventions

**ENTRYPOINT/CMD**

### DONT'S



- ▶ Avoid the "shell" form

### DO'S



- ▶ Use the "exec" form
  - ▶ Shell expansion will **not** happen!
- ▶ Use ENTRYPOINT and CMD together
- ▶ Use a "entrypoint-script"
  - ▶ Always "exec" (or "gosu")

**EXPOSE**



# DON'TS



- ▶ Avoid "docker run -P"

### DO'S



- ▶ Do document the ports your application needs exposed

ENV

## DO'S



- ▶ Use them for documentation/refactoring
- ▶ Use `docker run <image-name> env`
  - ▶ Or `docker inspect`
- ▶ Use them in conjunction with ARG
  - ▶ And default values

```
ARG PROJECT_VERSION
```

```
ENV PROJECT_VERSION ${PROJECT_VERSION:-2.3}
```

```
# ENVs are available in the container
```

```
# docker run -itP --rm <image-name> env
```

```
# RUN supports ENV expansion
```

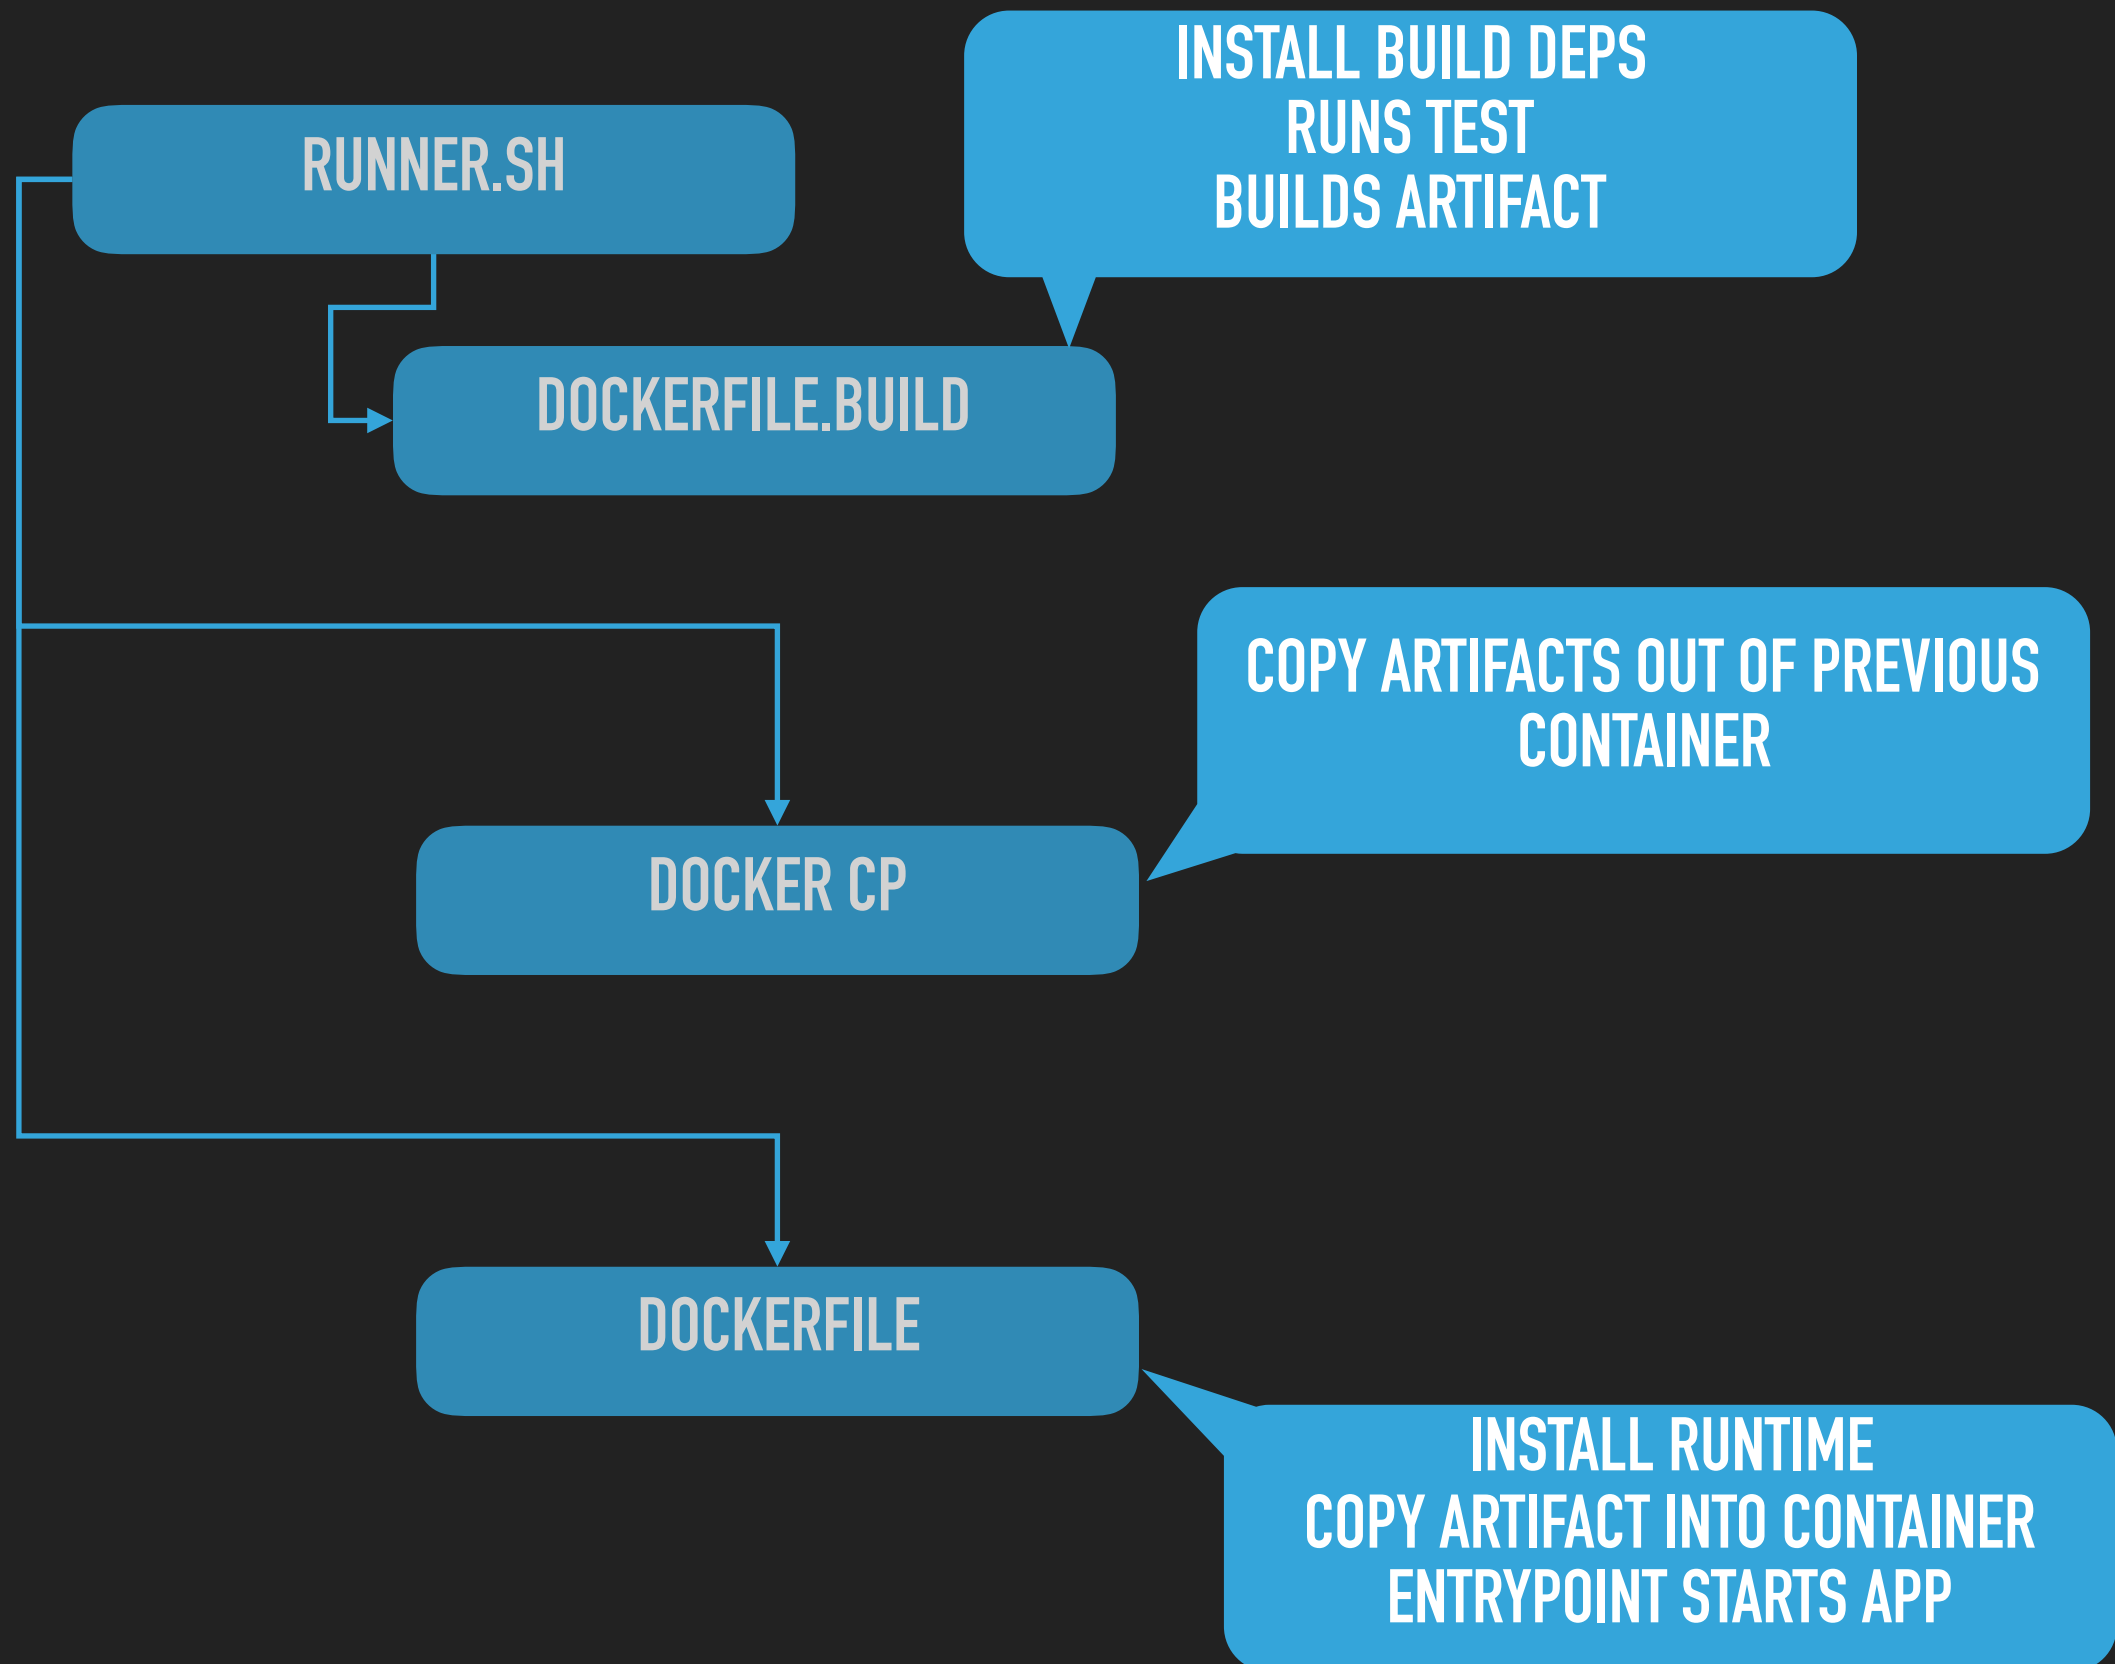
```
# RUN wget -O ${PROJECT_URL}
```

# CLOSING THOUGHTS

### DO'S



- ▶ Do create your own "ancestry" of images
- ▶ Do create *multiple* Dockerfiles per project (if necessary)
  - ▶ Dockerfile.build (installs JDK, build dependencies, builds war)
  - ▶ Dockerfile (installs JRE, and runtime e.g tomcat) with artifact
  - ▶ A wrapper script to tie it altogether





ALL DONE

---

**THANKS!!**

**RAJU GANDHI**

## RESOURCES

---

- ▶ [Best practices for writing Dockerfiles](#)
- ▶ [Guidance for Docker Image Authors](#)
- ▶ Dockerfiles for reference
  - ▶ [redis](#)
  - ▶ [Jenkins](#)
  - ▶ [Postgresql](#)